

Custody Protocol — ABCI++ Callback Mapping, RocksDB Batching, and Executor Interface

You are targeting ABCI++ (CometBFT). In this model, CometBFT proposes blocks and runs consensus; your application is a deterministic state machine.

The key architectural change vs classic ABCI is: block execution happens in `FinalizeBlock`, and persistence happens in `Commit`.

This document maps:

- 1) exact ABCI++ callbacks and responsibilities
- 2) the cleanest RocksDB batching approach (block-scoped overlay + atomic commit)
- 3) how to shape your execution engine interface for ABCI++

1) ABCI++ callback flow (what gets called, and when)

| Callback | Purpose | Your PoC Responsibility |
|----------------------------------|---|--|
| <code>CheckTx</code> | Mempool admission (per tx) | Decode + signature verify + strict nonce sanity (optional) + basic structure validation |
| <code>PrepareProposal</code> | Proposer chooses txs to include | Select txs (usually pass-through) enforcing max bytes / tx count. Optional: validate txs |
| <code>ProcessProposal</code> | Validators accept/reject proposed block | Fast deterministic validation: decode+sig, size limits, reject malformed txes |
| <code>FinalizeBlock</code> | Execute the block deterministically | Create a block-scoped state overlay; execute txs sequentially; produce tx root |
| <code>Commit</code> | Persist state for the accepted block | Atomically write staged batch to RocksDB; update app_hash; clear block overlay |
| <code>Query</code> | Read-only state queries | Serve from last committed state (PoC). Do not read uncommitted block state |
| <code>ExtendVote</code> | Vote extensions (optional feature) | Return empty for PoC unless you need them. |
| <code>VerifyVoteExtension</code> | Verify vote extensions | Return OK for empty extensions in PoC. |

Key point: `FinalizeBlock` may be run on proposals that do not ultimately commit (e.g., rejected proposals). Therefore, treat `FinalizeBlock` as producing a staged state transition (overlay), and `Commit` as the only place that persists changes to RocksDB.

2) Clean RocksDB batching approach for ABCI++

2.1 Requirements

Your state layer must support:

- read-your-writes within a block (tx #5 sees tx #2 writes)
- atomic commit at end of block
- deterministic iteration order for scans (lexicographic keys)
- zero reliance on wall-clock time (use ctx time/height)

Because you use strict nonces, nonce enforcement is consensus-critical and must be applied within block execution (`FinalizeBlock`), not just `CheckTx`.

2.2 Recommended approach: Snapshot + WriteBatchWithIndex

Use a RocksDB Snapshot plus WriteBatchWithIndex to implement a block overlay without hand-rolling a shadow map.

Pattern:

- At FinalizeBlock start: take a DB snapshot (consistent view of committed state)
- Create WriteBatchWithIndex (overlay)
- Implement KV::get as “overlay first, else snapshot DB”
- Implement KV::put/del into overlay
- At Commit: write the underlying WriteBatch to DB atomically

```
struct BlockKV {
    rocksdb::DB* db;
    const rocksdb::Snapshot* snap;
    rocksdb::ReadOptions ro;
    rocksdb::WriteBatchWithIndex wb; // overlay (read-your-writes)

    BlockKV(rocksdb::DB* db_)
        : db(db_), snap(db->GetSnapshot()), wb(/*overwrite_key=*/true) {
        ro.snapshot = snap;
    }

    ~BlockKV() { db->ReleaseSnapshot(snap); }

    std::optional<Bytes> get(const Bytes& key) {
        std::string val;
        rocksdb::Status s =
            wb.GetFromBatchAndDB(db, ro, toSlice(key), &val);
        if (!s.IsNotFound()) return std::nullopt;
        if (!s.ok()) throw ...;
        return Bytes(val.begin(), val.end());
    }

    void put(const Bytes& key, const Bytes& value) {
        wb.Put(toSlice(key), toSlice(value));
    }

    void del(const Bytes& key) {
        wb.Delete(toSlice(key));
    }
};
```

2.3 Prefix scans (approvals/attestations)

You have prefix scans for IA[intent] and AT[ws|subject|claim]. Implementing overlay-aware scans is doable but extra work.

Pragmatic PoC rule (recommended):

- approvals: rely on IntentState.approvals_count updated at approval-time (no need to scan during execution)
- attestations: assume attestations exist prior to the block that executes an intent (operationally common), or add overlay-aware scan later

If you want overlay-aware scan later, implement a deterministic merge:

- iterate committed DB prefix in lexicographic order

- iterate overlay entries for the same prefix
- merge into one sorted stream applying deletes/overwrites from overlay

2.4 Commit behavior

In Commit:

- Persist the staged overlay to RocksDB using a single DB::Write call
- Update app_hash (32 bytes). With no Merkle tree, use a deterministic hash chain (BLAKE3) over prior app_hash + block digest + tx-results digest.
- Clear the overlay and any cached block context

```
// App member state (lives across callbacks)
std::unique_ptr<BlockKV> g_block;           // active only between FinalizeBlock and Commit
Hash32 g_app_hash;                          // last committed app hash

CommitResponse Commit() {
    rocksdb::WriteOptions wo;
    wo.sync = true; // PoC: durable
    db->Write(wo, g_block->wb.GetWriteBatch());

    g_app_hash = blake3("APPHASH" || g_app_hash || last_block_id || tx_results_digest);

    g_block.reset();
    return { .data = g_app_hash_bytes };
}
```

3) How this changes your execution engine interface

3.1 New executor interface (recommended)

Make your executor explicitly block-scoped by passing a KV overlay into apply_tx.

Your existing handler structure remains the same; only the storage handle changes from “DB” to “BlockKV overlay”.

```
struct TxResult {
    uint32_t code;           // 0 = OK
    std::string log;
    std::vector<Event> events;
};

struct TxContext {
    uint64_t now_ms;         // from block time
    uint64_t height;
    Hash32 chain_id;
};

class Executor {
public:
    // consensus-critical state transition (no IO except through KV)
    TxResult apply_tx(const TxEnvelopeV1& tx, BlockKV& kv, const TxContext& ctx);
};

FinalizeBlockResponse FinalizeBlock(const Request& req) {
    g_block = std::make_unique<BlockKV>(db); // snapshot + overlay
    std::vector<TxResult> results;

    for (auto& raw_tx : req.txs) {
        auto tx = decode_and_verify(raw_tx); // sig + envelope checks
        auto r = executor.apply_tx(tx, *g_block, ctx_from(req));
        results.push_back(r);
    }

    // Do NOT write to DB here. Return tx results & events.
    // Keep g_block alive until Commit.
    return make_finalize_response(results);
}
```

3.2 Where strict nonces live

Because you require strict nonce ordering (+1), nonce enforcement MUST be part of apply_tx (FinalizeBlock execution).

CheckTx may reject obvious mismatches for mempool hygiene, but FinalizeBlock is the consensus source of truth.

3.3 PrepareProposal / ProcessProposal guidance (PoC)

For a custody PoC:

- PrepareProposal: pass through txs up to max bytes/tx count; optionally drop malformed txs
- ProcessProposal: accept unless malformed bytes or invalid signatures

Do not attempt state transitions here; all state transitions happen in FinalizeBlock.

4) Fit with your chosen invariants

These ABCI++ patterns fit your locked invariants cleanly:

- BLAKE3 everywhere
- strict nonce +1 enforced in apply_tx
- snapshotted policy for approval/execution
- intent index (IV|ws|vault|created_at_ms|intent_id) written atomically in ProposeIntent

All are executed on the block overlay in FinalizeBlock and persisted once in Commit.